


Agent Dispatcher, before and now

What changed, how the pieces connect, and where savings may come from.



BEFORE THIS UPDATE

Role-led setup

Simple work could still carry optional role/reporting overhead.

Excerpt-only estimate

Guidance, map data and metadata sat outside that budget.

Fact map + text retrieval

Manual cache updates; no persistent structural graph.

Repeated source passages

Another preparation could resend unchanged excerpts.

IMPLEMENTED NOW

Three work profiles

Direct, Guided and Coordinated match setup to the task.

One compact-package budget

Supplied guidance, excerpts, maps and metadata count together.

Maintained indexes + task graph

Related code helps decide which source passages to include.


Optional retained-evidence reuse

Exact references replace passages still in the same context.


Same host model and permissions. New preparation controls. Cost and speed gains are not yet measured.

The request now chooses its own level of setup

The profile is chosen from the request, without an extra router or model call.



Small tasks could already skip the context helper; the new direct path also removes optional surrounding setup.



Forced roles and requested workflows retain their guidance. Security, ambiguity and multi-file work stay guided.

One budget for the whole preparation packet

The change expands what is counted; it does not claim a specific reduction in model usage.

BEFORE: EXCERPTS COUNTED

SOURCE EXCERPTS

Their estimated size was bounded.

Role / guide instructions

Map facts + resource locations

Diagnostics + other metadata

Those additions were outside the excerpt-only accounting boundary.

NOW: ALL RETURNED MATERIAL COUNTS

Supplied role + selected guides

Selected source excerpts

Fact map + task graph


Metadata + diagnostics + exclusions

ONE SERIALIZED JSON PACKET

DEFAULT TOTAL LIMITS	4,000	8,000	18,000	4 chars/token estimate. Host instructions, history and later reads are outside this packet.
Estimated tokens	Small	Standard	Complex	

One source scan feeds two persistent indexes

Automatic maintenance runs during substantial preparation. It is not a background watcher.



FIRST USE
Build owned local indexes

UNCHANGED
Revalidate; skip identical writes


SOURCE CHANGED
Re-derive and refresh safely

PARTIAL / EXCLUDED
Use allowed evidence; defer saving

Every preparation still checks current sources and derives the graph. Persistence helps retain structure; it does not remove scan or parsing work.

Find related code through relationships

Illustrative Python example: fix validate_session. These are static code links, not a runtime trace.




Task-seeded PageRank ranks nearby relationships. Debugger, reviewer and tester roles give test files extra weight.

Every relationship includes a source line, source fingerprint, extraction method and resolved / inferred label.

Send unchanged passages once, when context is retained

Reuse is optional and explicit. Each preparation still checks the current source files.



Private ledger outside the project

Stores fingerprints, not source text or task prose. It does not establish cross-chat memory.

Successful output first, ledger update second

The CLI writes and flushes the packet before recording delivery. Unsafe state falls back to full evidence.

What happens when you ask for a multi-file fix

Example request: "Fix the session-validation bug and check its callers."

- 1

Choose Guided

Select the relevant role. Scope and risk justify prepared investigation.
- 2

Read permitted sources once

Fingerprint and redact files; respect no-read exclusions before content reads.
- 3

Use structure and relevant changes

Maintain local indexes. Related callers, dependencies and tests influence ranking.
- 4

Build the compact packet

Supply role text and useful evidence. Apply optional reuse and the total budget.
- 5

Work in the host and verify

Claude/Codex uses native tools and required checks; useful independent work may be delegated.

WHY IT CAN HELP

Less setup

Simple tasks use Direct instead.

Less rediscovery

A related caller can be surfaced with its source evidence.

Less repetition

Retained passages can be referenced rather than resent.

More honest limits

Missing or inferred relationships remain visible.

Changed focus or relevant source? Prepare again. Test commands in a map are suggestions, not evidence of a test pass.

What was implemented, and what remains to prove

The preparation layer was updated in source and generated packages for both supported hosts.

IMPLEMENTATION MAP

Work profiles

SKILL.template.md + Codex template

Selection + changed-file focus

context.py

Whole-packet budget

context_packet.py

Retained-evidence ledger

context_reuse.py

Automatic fact maintenance

project_map.py

Structural extraction + ranking

project_graph.py

Host packaging + activation

build.py / build_codex.py + adapters

Future extensions: broader Tree-sitter/LSP parsing, semantic search, git co-change history, coverage-backed tests and framework-specific flows.

14 offline suites passed

At handoff: packaging, budgets, freshness, exclusions, graph retrieval and reuse delivery checks.

Bounded first version

Python symbols/calls; JS/TS import hints. Index: 80 sources, 240 nodes, 400 edges. Task view: up to 12 nodes / 16 edges.

Next: measure the outcome

Use matched tasks, model and effort to compare correctness, time, token use and interventions.

No cost or speed advantage has been established by these offline passes.