


Three workflows. The same host capabilities.


Dispatcher runs inside Claude Code or Codex. It adds a curated preparation and verification process; the host still runs the model and tools.



Already native: tools, skills, subagents and context management. Dispatcher's opportunity is better selection, less repeated work, and more consistent evidence - measured over the whole task.

What happens inside our Dispatcher

The current working tree combines model judgment, local Python helpers, and host execution. It does not run a separate model service by default.



Version note: this diagram describes the current repository, including uncommitted verification/preferences work. The installed Codex pack is older. The role catalog contains 27 roles and 79 guides; they are not all loaded for each task.

Our small Claude pilot shows overhead to remove

Five frozen fixture tasks, each repeated twice per condition: 10 stock runs and 10 Dispatcher runs. Claude Opus 5, high effort, subscription sign-in.


10 / 10

STOCK

10 / 10

DISPATCHER

Both passed automated artifact and owned-directory scope checks.
These checks do not establish all-around quality or claim accuracy.



Paired median overhead: +15.6 seconds, +1,486.5 output tokens, +\$0.158 reported estimate per matched pair. Paired medians differ from subtracting the separate condition medians above.

Interpretation: this pilot does not demonstrate a cost or quality advantage. It used a frozen package, not every current pending change. The dollar values are provider-reported estimates, not subscription invoices. There is no Codex result in this comparison.

The saved trace audit also found some unsupported or overstated reporting. Human claim-accuracy and intervention ratings are absent here; a separate older 60-run pilot still has 24 required reviews pending.

Six changes worth testing for lower total usage

Each proposal targets a specific source of overhead. Priority reflects the inspected implementation and pilot; savings remain unmeasured.

01

FIRST

Make the fast path real

For simple tasks, skip extra role and helper loading when native behavior suffices. Strengthen the existing bypass rules.

Measure: preparation tokens and time on small tasks.

02

FIRST

Budget the whole context packet

Count guidance, metadata and diagnostics as well as source excerpts. Keep normal output compact, with details on request.

Measure: total packet size plus task acceptance.

03

NEXT

Reuse only unchanged evidence

Track already-read excerpts and guide versions. Send changed passages; preserve source links and invalidate stale facts.

Measure: repeated reads and total input usage.

04

CONDITIONAL

Match model and effort to the job

Use an explicitly configured economical profile for clear work. Escalate when objective checks fail or uncertainty remains.

Measure: complete task cost, including escalations.

05

NEXT

Budget delegation

Prefer one agent for tightly coupled work. Give useful workers narrow inputs and bounded result formats; avoid duplicating searches.

Measure: all child usage, coordination and elapsed time.

06

NEXT


Retrieve around changed code

Seed retrieval from changed lines and symbols, then inspect callers and relevant tests. Expand only when evidence is missing.

Measure: needed-file recall versus returned bytes.

Spend extra effort where it changes the outcome

Use observable failure and task risk to choose the next step. Stronger reporting and targeted verification can reduce both mistakes and rework.



Optimize cost per accepted task

The comparison must count all attempts, child agents and verification work. Lower token use is useful only when the completed result still meets the task.

A Stock host

Same task, model, effort, tools, project rules and source snapshot.

B Current Dispatcher

Freeze the package and run its normal preparation and evidence workflow.

C One proposed change

Change one feature at a time. Test model routing separately from prompt/context changes.

Measure the whole job

Outcome Acceptance, correctness, scope and unsupported claims.

Usage Input, cache reads/writes, output, tools and all children.

Time Elapsed time, retries and avoidable human corrections.

Money API bills when available; reported estimates labeled separately.

Use the existing evaluation runner

1. Add representative bugs, features, reviews and documentation tasks.
2. Randomize paired runs; keep every attempt and review results blind where feasible.
3. Report Claude and Codex separately. Track cache conditions and freeze versions.
4. Adopt a change when quality is maintained at lower cost, or the quality gain justifies its cost.

Suggested build order: compact task contracts and helper output; evidence-backed reporting; then selective reuse, measured routing and task recovery. Run evaluation on demand. Keep task-observer disabled.

What this report is based on

Official host documentation plus the local repository and saved evaluation artifacts inspected on 20 September 2026. Web-source links are clickable.

HOST CAPABILITIES

- C1 Claude Code: how it works
- C2 Claude Code: skills
- C3 Claude Code: subagents
- C4 Claude Code: costs and usage
- C5 Claude Code: memory
- O1 Codex: customization and progressive loading
- O2 Codex: subagents and model/effort controls
- O3 Codex: pricing and allowance usage
- F1 Firstmate: recovery and supervision architecture

Terms used in this report

- Host:** Claude Code or Codex, which runs the model and tools.
- Role:** a focused set of working instructions.
- Context:** information supplied to the model.
- Receipt:** a record of checks actually executed.
- Accepted task:** a result that passes the task's agreed checks.

LOCAL IMPLEMENTATION AND MEASUREMENTS

- R1 **Routing and delegation**
adapters/codex/SKILL.template.md; DELEGATION.template.md
- R2 **Context selection and project facts**
CONTEXT.template.md; context.py; PROJECT-MAP.template.md; project_map.py
- R3 **Pending check receipts and preferences**
VERIFICATION.template.md; verification.py; preferences.py
- E1 **Focused pilot: report.md and batch.json**
dist/evals/opus5-reliability-fix-20260920/
batches/20260920T053330Z-46a6325f/
- E2 **Trace audit and limitations**
dist/evals/opus5-reliability-fix-20260920/
SUMMARY.md; analysis/independent-trace-audit.md
- E3 **Earlier pilot, 24 reviews pending**
dist/evals/opus5-20260919/
corrected-pilot/report.md
- E4 **Evaluation method and accounting caveats**
evals/end_to_end/README.md

Scope: current working-tree design, including pending changes; installed Codex pack differs. Host features depend on version, configuration and account. The pilot used Claude Code 2.1.150 and a frozen package (fd75da1af004); its results cannot establish current-tree or Codex performance.